

CS-233: Introduction to Machine Learning

Course Project – Milestone 2

Welcome to the second Milestone of the project for the course Introduction to Machine Learning. Similar to Milestone1, our goal will be to implement a few of the methods seen in the course and to test them on a dataset. We will be evaluating them on evaluation metrics specific for the tasks and present our results with a project report at the end of each milestone.

As deliverable, you are expected to fill a code framework: python files `.py` that implement these specific methods, and then a main python script `main.py` to call them on the data. We will use a subset of the [Gaming and Mental Health Dataset](#) dataset, which is dataset exploring gaming habits and mental health: the goal is to predict addiction level from gaming habits.

You are encouraged to complete the project as you encounter the methods in the course. You should also come to the coding exercise sessions where you can discuss the project with the assistants.

You can find more information in the respective chapters below.

Important Information:

- This will be a group project and must be done in **groups of 3**.
- You may post your questions about the project on the Ed.
- *You must implement all the code by yourself.* You are not allowed to import and use external libraries (e.g., scikit-learn), unless it is *explicitly* said otherwise.
 - NumPy, and matplotlib are fine to use.
- **Deadline:**
 - **Milestone 2: 29.05.2026 (17:00)**

Grading

Each milestone is worth 10% of your final grade and will be graded as following: (**10pt**)

- Code (**3pt**): the framework code was filled as asked and works as expected.
- Clarity and presentation (**3pt**): the report is clear, well presented and structured.
- Discussion (**4pt**): discussion of the work, what you did, what worked and didn't, how you made your choices, reporting of the results, comparison of methods, etc.
- *Creativity* (**1pt**): if you went beyond what is asked, e.g., by implementing something more, an interesting analysis or experiment, an especially well written code, etc.
 - Note: while you can earn up to 11pt, each milestone will be capped at 10pt.

Respect of the page limit: We will remove 1pt if the report does not respect the page limit.

The Framework

Structure

The framework is organized with the following folder structure: (you can download it from Moodle)

```
<sciper1>_<sciper2>_<sciper3>_project/  
  main.py  
  test_ms2.py  
  report.pdf  
  data/  
  src/  
    __init__.py  
    utils.py  
    activations.py  
    losses.py  
    methods/  
      __init__.py  
      dummy_methods.py  
      mlp.py  
      kmeans.py
```

- Please do not rename these files, nor move them around.
- You are **allowed** to add additional python files.
- Your report will also be a part of this folder structure (as `report.pdf`).

Please include it in your framework when you are preparing your submission. You will be submitting the zipped "project" folder.

- Download the data from moodle (features.npz) and place it into the `data/` folder, *but make sure you remove it for the submission*, otherwise the file size will get very large!
- Make sure to take a look into `src/utils.py` for useful functions you could use.

Implementation

Some parts of the code are provided to help you get started. It will be your task to fill in the missing parts to get all the methods running on your data. You have to 1) fill the python files implementing the methods, the losses and the activations, and 2) write the main script that runs these methods on the dataset.

- **Methods:** The methods are all implemented in the form of Python classes, under `src/methods/`. We have provided the implementation of a "dummy" classifier that returns random predictions as an example, in `src/methods/dummy_methods.py`. Study this class well, you can use it as templates to code your own methods. It contains some essential functions:
 - `__init__(self, ...)`
 - ♣ This function is used to initialize an object of the class. You can give it arguments that will be saved in the object, such as the value of hyperparameters.
 - `fit(self, training_data, training_labels)`
 - ♣ This is the training or fitting step of the model. Training data with corresponding labels are used to estimate the parameters of the model. This method should be called before the `predict` function.
 - `predict(self, test_data)`
 - ♣ It generates predictions for the given data. It should be called after the `fit` function (you cannot predict without a trained model!).
- **Main script:** Contrary to the notebooks of the exercise sessions, we run python as a script in the terminal (see below for details on how to do it). You need to complete this script `main.py` to apply the methods to your data and evaluate them. Some pointers are given in the form of comments in the file. In brief, you are expected to follow these general steps:
 - Loading the data
 - Preparing / pre-processing the data
 - Initialize the selected method
 - Training and evaluating the selected method
 - Reporting the results
- This script can be broken down in two main parts:
 - The main function `def main(args):` which is where we write the code we run.

- o And some code starting with `if __name__ == '__main__':`. This is some pythonic way of coding “run the following instructions only if this file is a script”. This is where we parse the arguments that were given in the command line and then call the `main()` function.

Running the Project

You will be running the Python scripts from your terminal. For example the main script can be launched from inside the project folder as:

```
python main.py --data_path <where you placed the data folder> --method
dummy_classifier --task classification
```

You can specify the method and many other arguments from the terminal by writing them as `--arg value` in the command above. The arguments are defined at the bottom of the `main.py` script. We encourage you to check how they work, what their default values are, and you can even add your own!

Test Script

We also provide you with a test script `test_ms2.py`. This script is for you to verify your project folder and code structure, as well as testing your method on some *very easy* cases.

- These easy cases are by far not exhaustive and you should verify the correctness of your code by your own means as well!
- The script can be launched from inside the project folder as
 - o `python test_ms2.py`
- By default, it hides any `print()` from your code. You can optionally pass it the argument `--no-hide` to re-enable the printing.

Sample results

We do not provide sample results for this milestone, as you should already have an idea of what is possible on this dataset from the first milestone.

The Gaming and Mental Health Dataset

The Gaming and Mental Health Dataset is synthetic dataset of samples collected from 10 million users. However we provide you a downsampled dataset containing 2000 samples, and featuring 15 columns.

13 of these are the features you will use to predict either the addiction score (regression task) or the addiction level (classification task).

- The dataset is split into 1600 train samples and 400 test samples
 - The regression labels (addiction score) are integers between 0 and 10.
 - The classification labels (addiction level) belong to {Low, Medium, High} and encoded as integers {0, 1, 2} respectively.
- You should download the dataset from moodle and place into the `data` folder. `main.py` loads the dataset as 6 numpy arrays:
 - `train_features` of shape (1600, 13)
 - `test_features` of shape (400, 13)
 - `train_labels_reg` of shape (1600,)
 - `test_labels_reg` of shape (400,)
 - `train_labels_classif` of shape (1600,)
 - `test_labels_classif` of shape (400,)

The Tasks and Metrics

Milestone 2

For Milestone 2 (MS2), you are expected to:

- Code from scratch a **Multi-Layer Perceptron (MLP)** and use it to perform both **Regression** (predicting the **addiction score**) and **Classification** (predicting the **addiction level**):
 - You have to code by hand the forward pass, back propagation and weight updates. You are not allowed to use `torch.nn`.
 - You should also complete the losses and activation functions (including gradient calculation to be used for training) that can be found in `losses.py` found `activations.py`. You are free to code additional losses and activation functions.

- Code from scratch **K-Means**, to learn K clusters in an unsupervised manner. Although K-Means is an unsupervised clustering algorithm, in this project we expect you to use it for classification in the following way:
 - Assign a class label to a cluster by *voting*: choosing the most represented labels among the samples associated to the cluster.
 - Predict the class of an unseen data sample, by assigning it to a cluster, and outputting that cluster's label as the prediction.
- The main script that calls them.

And you will need to write a **2-pages report** about what you did.

- **Deliverable:** You need to submit **your code** and a **2-pages report** (zipped together). The project report should be part of your folder structure (see above!)
 - You should name the project folder `<sciper1>_<sciper2>_<sciper3>_project`, and the zipped file should have the name `<sciper1>_<sciper2>_<sciper3>_project.zip`. Please make sure that when you unzip it, it extracts the folder `<sciper1>_<sciper2>_<sciper3>_project`.
- **Code:** You will need to complete the relevant parts of the following python files for this milestone. Please look inside the files for more information.
 - `main.py`
 - ♣ The main python script we use to run our methods. While the data is loaded for you, you should prepare it correctly (e.g., normalization, bias, validation set,...), and then call the methods to train and predict.
 - `src/methods/mlp.py`
 - ♣ Implement a multi-layer perceptron, as seen in the class. Note that we will be implementing everything by hand, including back-prop.
 - `src/methods/kmeans.py`
 - ♣ Implement K-means clustering algorithm and its extension to classification.
 - `src/losses.py`
 - ♣ Implement Mean-Square Error (MSE) and Cross-Entropy losses. You can code additional losses.
 - `src/activations.py`
 - ♣ Implement Sigmoid and ReLU activation functions. You can code additional activation functions.
- **Report:** You need to write a concise *2-pages* report about your work. The project report should include your methodology and a brief summary of the results you have achieved with the methods.
 - We advise to follow a structure like "Introduction, Method, Experiment/Results, and Discussion/Conclusion".

- Points to consider when writing your report: your data preparation, the methods you tried, their performances, how you selected the hyperparameters, your MLP's architecture and how you selected it, how do the *final* models compare on the test data, etc.
 - You should explain what hyperparameters you tried for these methods and what results you achieved. For instance, you can show simple visualizations of your validation performances with respect to the different hyperparameters.
 - Stay concise and respect the page limit. You can use double column if needed.
- **Test:** Make sure that `test_ms2.py` runs without any problems! This ensures that your code compiles and your implementations can be imported.